

VIENNA UNIVERSITY OF TECHNOLOGY

360.252 COMPUTATIONAL SCIENCE ON MANY CORE ARCHITECTURES

INSTITUTE FOR MICROELECTRONICS

Exercise 8

Authors:

Camilo TELLO FACHIN
12127084

Supervisor:

Dipl.-Ing. Dr.techn. Karl RUPP

December 13, 2022



TECHNISCHE
UNIVERSITÄT
WIEN

Vienna University of Technology

Abstract

Here documented the results of Exercise 8.

Contents

1	Algorithm in CUDA with double precision arithmetic and data types (3/4 Points)	1
1.1	OpenCL Kernel to Compute Dot Product of two vectors	1
1.2	Compare Performance of OpenCL Kernels to Cuda Kernels	2
1.3	Compare Performance of OpenCL Kernels on K40 CPU	3
1.4	Measure Time it takes to build M OpenCL Programs	4
2	Bonus: Implement CUDA+OpenCL (CUCL) Approach (0/1 Points)	5

1 Algorithm in CUDA with double precision arithmetic and data types (3/4 Points)

1.1 OpenCL Kernel to Compute Dot Product of two vectors

The given code skeleton `vector_add.cpp` was adapted in a very rudimentary way without shared memory or any other bamboozles! The Sum of the entries later computed on the CPU.

C++ Code Changes in `vector_add.cpp`

```
1  const char *my_opengl_program = R"(
2  #pragma OPENCL EXTENSION cl_khr_fp64 : enable // required to enable 'double' inside OpenCL programs
3
4  __kernel void vec_add( __global double *x,
5                        __global double *y,
6                        unsigned int N)
7  {
8      for (unsigned int i = get_global_id(0);
9           i < N;
10          i += get_global_size(0))
11          x[i] = x[i] * y[i];
12  })";
13
14
15  double dot_product = 0;
16  for(int i = 0; i < vector_size; i++){
17      dot_product += x[i];
18  }
```

One "advantage" of this blatantly simple kernel for the dot product, is that without introducing shared memory, partial exercise 3 works without problems if the device is set to CPU.

1.2 Compare Performance of OpenCL Kernels to Cuda Kernels

I used the CUDA Kernel for the dot product from exercise sheet 2 for comparison. For this rudimentary Kernel, the CUDA kernels outperform the OpenCL Kernels on both GPU's.

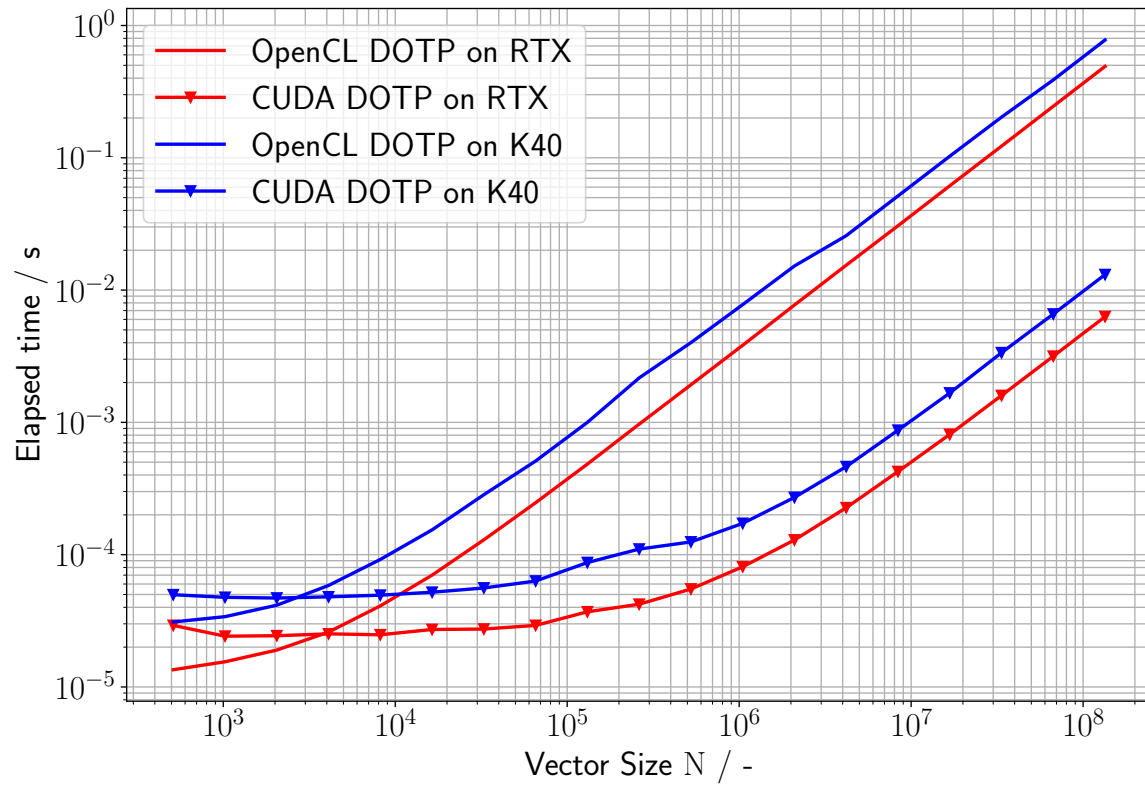


Figure 1: Kernel Performance Comparison

1.3 Compare Performance of OpenCL Kernels on K40 CPU

Like mentioned in part 1, here one simply has to change a 0 to a 1 and one can confirm in the output that the CPU is actually used. CUDA obviously still the fastest. OpenCL Kernel run on K40 is faster than on the CPU.




Figure 2: CPU GPU performance Comparison

1.4 Measure Time it takes to build M OpenCL Programs

not done :-).

2 Bonus: Implement CUDA+OpenCL (CUCL) Approach (0/1 Points)

not done :-).