

VIENNA UNIVERSITY OF TECHNOLOGY

360.252 COMPUTATIONAL SCIENCE ON MANY CORE ARCHITECTURES

INSTITUTE FOR MICROELECTRONICS

Exercise 9

Authors:

Camilo TELLO FACHIN
12127084

Supervisor:

Dipl.-Ing. Dr.techn. Karl RUPP

January 4, 2023



TECHNISCHE
UNIVERSITÄT
WIEN

Vienna University of Technology

Abstract

Here documented the results of Exercise 9.

Contents

1	Libraries (4/5 Points)	1
1.1	Boost.Compute (1/1 Points)	1
1.2	Thrust (1/1 Points)	1
1.3	VexCL (1/1 Points)	1
1.4	ViennaCL (1/1 Points)	1
1.5	Comparison to CUDA and OpenCL (0/1 Points)	1
2	HIP (Points 3/3)	2
3	BONUS: Relaxed Christmas Break (0/1 Point)	2
	Appendices	3
A	CPP CODE for Boost.Compute Library	3
B	CPP CODE for Thrust Library	5
C	CPP CODE for VexCL Library	6
D	CPP CODE for ViennaCL Library	7

1 Libraries (4/5 Points)

1.1 Boost.Compute (1/1 Points)

Successfully implemented it and it yields the correct dot product output. (See Appendix A)

1.2 Thurst (1/1 Points)

Successfully implemented it and it yields the correct dot product output. (See Appendix B)

1.3 VexCL (1/1 Points)

Successfully implemented it and it yields the correct dot product output. (See Appendix C)

1.4 ViennaCL (1/1 Points)

Successfully implemented it and it yields the correct dot product output. (See Appendix D)

1.5 Comparison to CUDA and OpenCL (0/1 Points)

Although It probably wouldve been interesting I just didnt have the time :-).

2 HIP (Points 3/3)

Successfully Implemented the classical CG Algorithm in HIP. Whereby "implementing" I mean replacing CUDA commands to HIP commands. A bit trickier where the Kernel calls, But eventually made it run and it still converges and yields corrects results (Relative residual error smaller than $1e-6$).

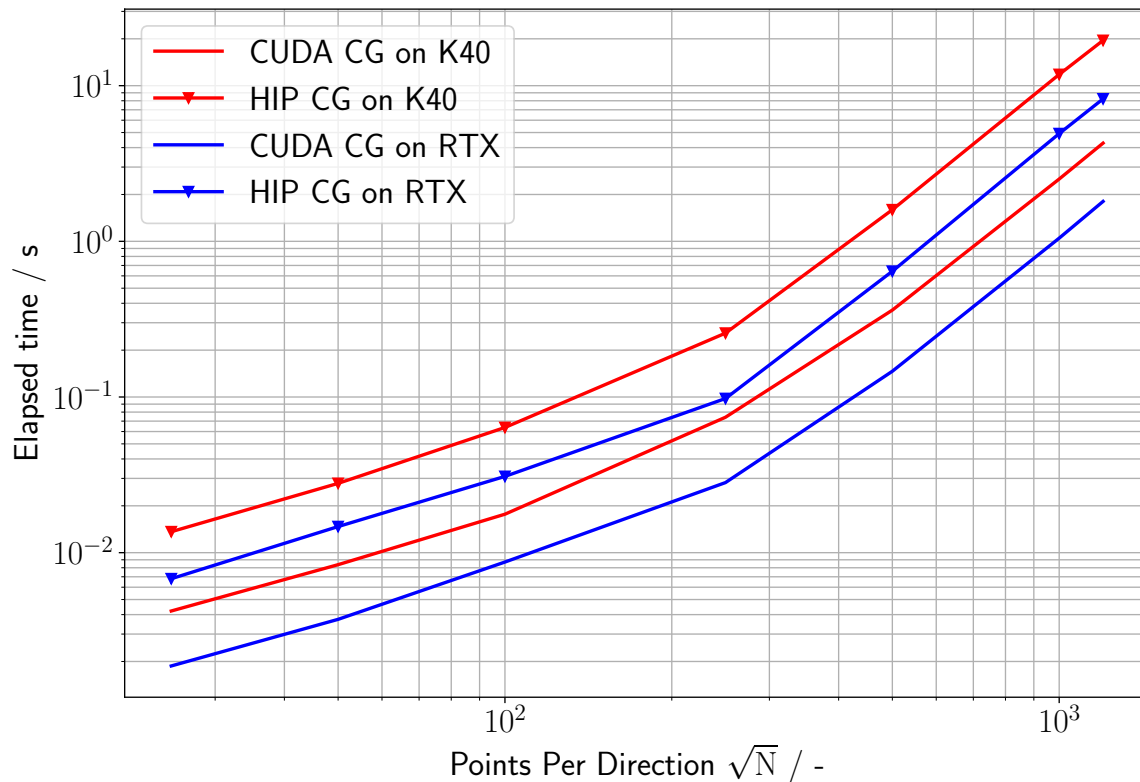


Figure 1: Comparison Conjugate Gradient Algorithm with CUDA and HIP

The HIP Implementation is slower for both GPU's!

3 BONUS: Relaxed Christmas Break (0/1 Point)

I had a very relaxed christmas break, therefore You could grant me the Bonus point! Jokes aside, I hope You had a very relaxed Christmas break! :-) Kind regards.

A CPP CODE for Boost.Compute Library

CPP CODE for Boost.Compute Library

```

1 namespace compute = boost::compute;
2
3 BOOST_COMPUTE_FUNCTION(float, boost_multiplication, (float x, float y),
4 {
5     return x * y;
6 });
7
8 BOOST_COMPUTE_FUNCTION(float, boost_addition, (float x, float y),
9 {
10     return x + y;
11 });
12
13 BOOST_COMPUTE_FUNCTION(float, boost_subtraction, (float x, float y),
14 {
15     return x - 2 * y;
16 });
17 int main(int argc, char *argv[])
18 {
19
20     int i, N, N_max, gentxt;
21     float dotP;
22     N_max = 100;
23     gentxt = 0;
24     for (i=1; i<argc&&argv[i][0]!='-'; i++) {
25         if (argv[i][1]=='N') i++, sscanf(argv[i], "%d", &N); // cmdline arg -N for adjusting max. count, if
                        none given N=100
26         if (argv[i][1]=='g') i++, sscanf(argv[i], "%d", &gentxt); // cmdline arg. -g for generating a txt, if
                        none given, no .txt NOT IMPLEMENTED
27     }
28
29     // get default device and setup context
30     compute::device device = compute::system::default_device();
31     compute::context context(device);
32     compute::command_queue queue(context, device);
33
34     // generate the 3 vectors: x=(1,1,...,1), y=(2,2,...,2) and dotp
35     std::vector<float> host_vector_x(N_max);
36     std::vector<float> host_vector_y(N_max);
37     std::vector<float> host_vector_prod(N_max);
38
39     std::fill(host_vector_x.begin(), host_vector_x.end(), 1.0);
40     std::fill(host_vector_y.begin(), host_vector_y.end(), 2.0);
41
42     // create the 3 vectors on the device
43     compute::vector<float> device_vector_x(host_vector_x.size(), context);
44     compute::vector<float> device_vector_y(host_vector_y.size(), context);
45     compute::vector<float> device_vector_prod(host_vector_prod.size(), context);

```

```
46
47 // transfer data from the host to the device
48 compute::copy( host_vector_x.begin(), host_vector_x.end(), device_vector_x.begin(), queue);
49 compute::copy( host_vector_y.begin(), host_vector_y.end(), device_vector_y.begin(), queue);
50
51 // calculate x+y within x
52 compute::transform(
53     device_vector_x.begin(),
54     device_vector_x.end(),
55     device_vector_y.begin(),
56     device_vector_x.begin(),
57     boost_addition,
58     queue
59 );
60
61 // calculate x-2*y within y
62 compute::transform(
63     device_vector_x.begin(),
64     device_vector_x.end(),
65     device_vector_y.begin(),
66     device_vector_y.begin(),
67     boost_subtraction,
68     queue
69 );
70 //after these two steps, vector x should have x_i + y_i in every entry, and vector y should have x_i -
    y_i in every entry and one can calculate the product in the third vector.
71
72 // calculate x_i * y_i
73 compute::transform(
74     device_vector_x.begin(),
75     device_vector_x.end(),
76     device_vector_y.begin(),
77     device_vector_prod.begin(),
78     boost_multiplication,
79     queue
80 );
81 //copy values back to the host
82 compute::copy(
83     device_vector_prod.begin(), device_vector_prod.end(), host_vector_prod.begin(), queue
84 );
85
86 dotP = std::accumulate( device_vector_prod.begin(), device_vector_prod.end(), 0);
87 std::cout << "<math>x+y, x-y> = " << dotP << " with N = " << N_max << std::endl;
88
89 return 0;
90 }
```

B CPP CODE for Thrust Library

CPP CODE for Thrust Library

```

1  int main(int argc, char *argv[]) {
2      int i, N, N_max, gentxt;
3      float dotP;
4      N_max = 100;
5      gentxt = 0;
6      for (i=1; i<argc&&argv[i][0]!='-'; i++) {
7          if (argv[i][1]=='N') i++, sscanf(argv[i], "%d",&N); // cmdline arg -N for adjusting max. count, if none
                        given N=100
8          if (argv[i][1]=='g') i++, sscanf(argv[i], "%d", &gentxt); // cmdline arg. -g for generating a txt, if
                        none given, no .txt NOT IMPLEMENTED
9      }
10
11     // generate vectors on host of length N_max
12     thrust::host_vector<float> h_x(N_max);
13     thrust::host_vector<float> h_y(N_max);
14     thrust::host_vector<float> h_prod1(N_max);
15     thrust::host_vector<float> h_prod2(N_max);
16     thrust::fill(h_x.begin(), h_x.end(), 1);
17     thrust::fill(h_y.begin(), h_y.end(), 2);
18     thrust::fill(h_prod1.begin(), h_prod1.end(), 0);
19     thrust::fill(h_prod2.begin(), h_prod2.end(), 0);
20
21     // transfer data to the device
22     thrust::device_vector<float> d_x = h_x;
23     thrust::device_vector<float> d_y = h_y;
24     thrust::device_vector<float> d_prod1 = h_prod1;
25     thrust::device_vector<float> d_prod2 = h_prod2;
26
27     thrust::transform(d_x.begin(), d_x.end(), d_y.begin(), d_prod1.begin(), thrust::plus<float>());
28     thrust::transform(d_x.begin(), d_x.end(), d_y.begin(), d_prod2.begin(), thrust::minus<float>());
29     thrust::transform(d_prod1.begin(), d_prod1.end(), d_prod2.begin(), d_prod1.begin(), thrust::multiplies<float>
        >());
30
31
32     // transfer data back to host
33     thrust::copy(d_prod1.begin(), d_prod1.end(), h_prod1.begin());
34     dotP = std::accumulate(d_prod1.begin(), d_prod1.end(), 0);
35
36     std::cout << "<x+y, x-y> = " << dotP << " with N = " << N_max << std::endl;
37
38     return 0;
39 }

```


C CPP CODE for VexCL Library

CPP CODE for VexCL Library

```

1  int main(int argc, char *argv[]) {
2      int i, N, N_max, gentxt;
3      double dotP;
4      N_max = 100;
5      gentxt = 0;
6      for (i=1; i<argc&&argv[i][0]!='-'; i++) {
7          if (argv[i][1]=='N') i++, sscanf(argv[i], "%d",&N); // cmdline arg -N for adjusting max. count, if none
                        given N=100
8          if (argv[i][1]=='g') i++, sscanf(argv[i], "%d", &gentxt); // cmdline arg. -g for generating a txt, if
                        none given, no .txt NOT IMPLEMENTED
9      }
10
11     vex::Context ctx(vex::Filter::GPU && vex::Filter::DoublePrecision);
12
13     std::cout << ctx << std::endl; // print list of selected devices
14
15     std::vector<double> x(N_max, 1.0), y(N_max, 2.0), prod(N_max, 0);
16
17     vex::vector<double> X(ctx, x);
18     vex::vector<double> Y(ctx, y);
19
20     vex::vector<double> F1 = X + Y;
21     vex::vector<double> F2 = X - Y;
22     vex::vector<double> P = F1 * F2;
23
24     dotP = std::accumulate(P.begin(), P.end(), 0);
25     std::cout << "<x+y, x-y> = " << dotP << " with N = " << N_max << std::endl;
26
27     return 0;
28 }

```

D CPP CODE for ViennaCL Library

CPP CODE for ViennaCL Library

```

1  int main(int argc, char *argv[]) {
2      int i, N, N_max, gentxt;
3      double dotP;
4      N_max = 100;
5      gentxt = 0;
6      for (i=1; i<argc&&argv[i][0]!='-'; i++) {
7          if (argv[i][1]=='N') i++, sscanf(argv[i], "%d",&N); // commandline arg -N for adjusting max. count, if none
                        given N=100
8          if (argv[i][1]=='g') i++, sscanf(argv[i], "%d", &gentxt); // commandline arg. -g for generating a txt, if
                        none given, no .txt NOT IMPLEMENTED
9      }
10
11     viennacl::vector<double> x = viennacl::scalar_vector<double>(N_max, 1.0);
12     viennacl::vector<double> y = viennacl::scalar_vector<double>(N_max, 2.0);
13
14     viennacl::vector<double> f1 = x + y;
15     viennacl::vector<double> f2 = x - y;
16     viennacl::vector<double> p = viennacl::linalg::element_prod(f1, f2);
17
18     dotP = std::accumulate(p.begin(), p.end(), 0);
19     std::cout << "<x+y, x-y> = " << dotP << " with N = " << N_max << std::endl;
20
21     return EXIT_SUCCESS;
22 }
```